

# Introduction to Database Systems - Indexes

Radim Bača

Department of Computer Science, FEECS

[radim.baca@vsb.cz](mailto:radim.baca@vsb.cz)

[dbedu.cs.vsb.cz](http://dbedu.cs.vsb.cz)

# Content

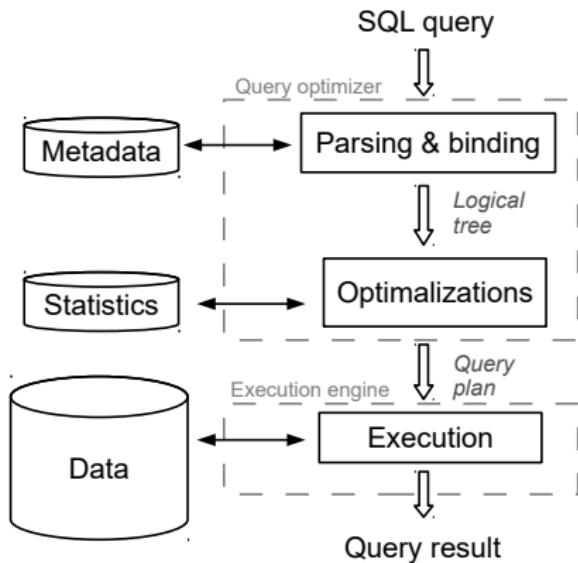
- Introduction
- What is an index?
- What is an appropriate index?

# Introduction

- Until now we were mostly working on a conceptual level of a database system:
  - Relational data model
  - SQL
  - Normal forms
- Today we will look closer under the hood of physical representation of the data

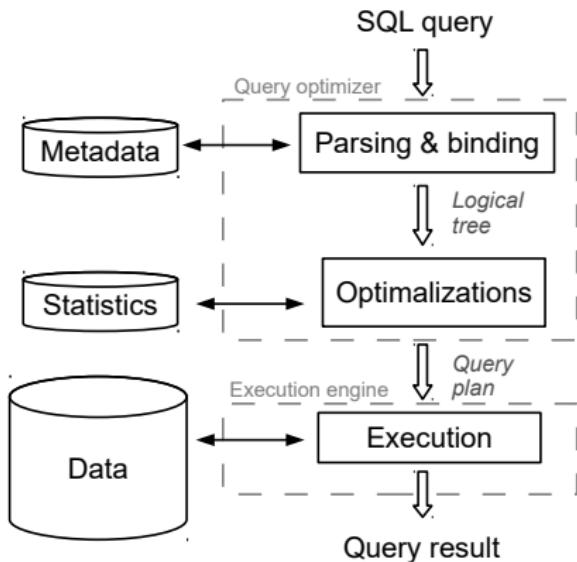


# SQL Query Compilation - Remainder



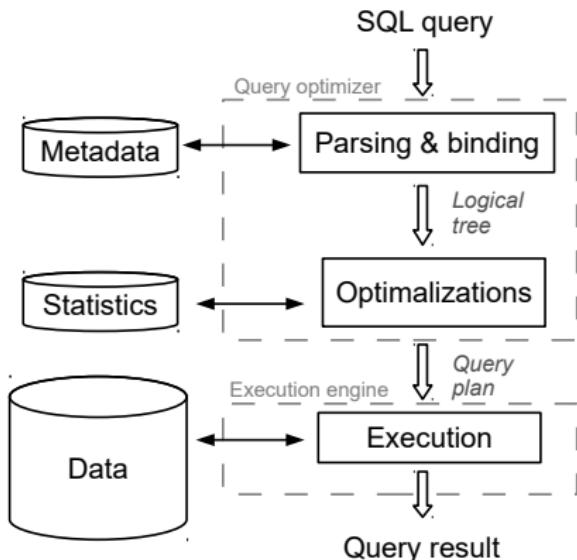
- SQL query optimization is a process which takes a SQL and output an query plan
- The process of plan creation should be deterministic
- If there are appropriate index then the optimizer uses it

# SQL Query Compilation - Remainder



- SQL query optimization is a process which takes a SQL and output an query plan
- The process of plan creation should be deterministic
- If there are appropriate index then the optimizer uses it

# SQL Query Compilation - Remainder

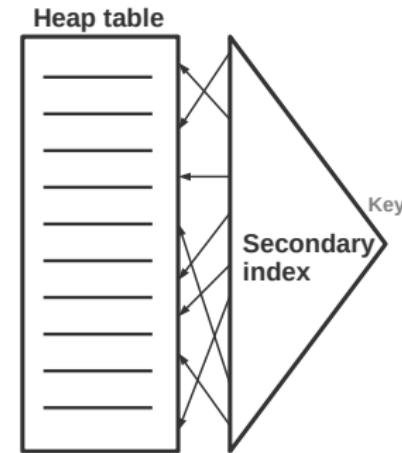


- SQL query optimization is a process which takes a SQL and output an query plan
- The process of plan creation should be deterministic
- If there are appropriate index then the optimizer uses it

But what is an index and what is an appropriate index?

# What is an Index?

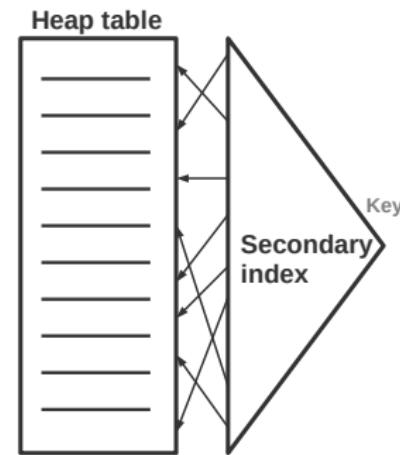
- Index allows fast search of a **value** according to a **key**
- It is data structure in a database system that can be created or dropped (see Lecture 6)
- The most typical data structure used in all database systems is a B+tree



## Why B+tree?

It has a following properties

- Search of an value according to a key in  $\log(n)$  time
- Reasonably fast of writes into the index (insert/update/delete)
- It support range scan in leaf nodes
- Data are organized into pages

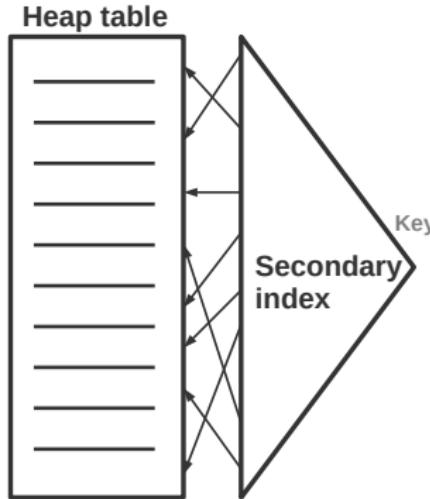


## Other Data Structures

- We may use other data structures if it is appropriate:
  - R\*tree - fast spatial search
  - Hash table - fast single record lookup
  - Skip list - lock free data structure
- It depends on availability in a database system
- However, `CREATE INDEX` creates a B+tree in most database systems

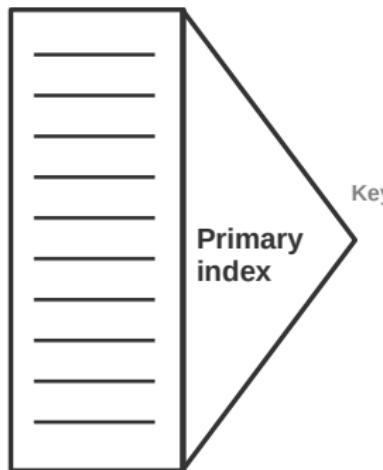
## Types of Indexes

- Secondary index - value is a pointer to a table



## Types of Indexes

- Primary index - value is a record of a table
- Database can contains just one such index



## Primary Index

- It has different names:
  - Clustered index (SQL Server, MySQL)
  - Index organized table (Oracle)
- It is not available in all database systems (PostgreSQL)

## When we Need an Index?

- Do we need an index for a table where only insert data and perform a query

SELECT \* FROM table?

- Definitely not
- Index helps us to process the predicates of a query more quickly
- The important observation:

The index is always connected to a SQL query (or a set of queries)

## When we Need an Index?

- Do we need an index for a table where only insert data and perform a query  
SELECT \* FROM table?
- Definitely not
- Index helps us to process the predicates of a query more quickly
- The important observation:

The index is always connected to a SQL query (or a set of queries)

## How to Create an Index

```
SELECT * FROM table WHERE atribut = hodnota
```



```
CREATE INDEX ix_table_attr ON table(atribut)
```

## How to Create an Index

```
SELECT * FROM table WHERE atribut = hodnota
```



```
CREATE INDEX ix_table_attr ON table(atribut)
```

- We create indexes according to the predicate in a query

## Example - SQL Server

- `SELECT * FROM Actor  
WHERE first_name = 'johnny'` →

Search using a sequential scan through the actor table

## Example - SQL Server

- `SELECT * FROM Actor  
WHERE first_name = 'johnny'`
- Now lets create an index ?
- `CREATE INDEX  
ix_actor_firstname  
ON Actor(first_name)` →

## Example - SQL Server

- `SELECT * FROM Actor  
WHERE first_name = 'johnny'`
- Now lets create an index
- `CREATE INDEX  
ix_actor_firstname  
ON Actor(first_name)`



The index is not used



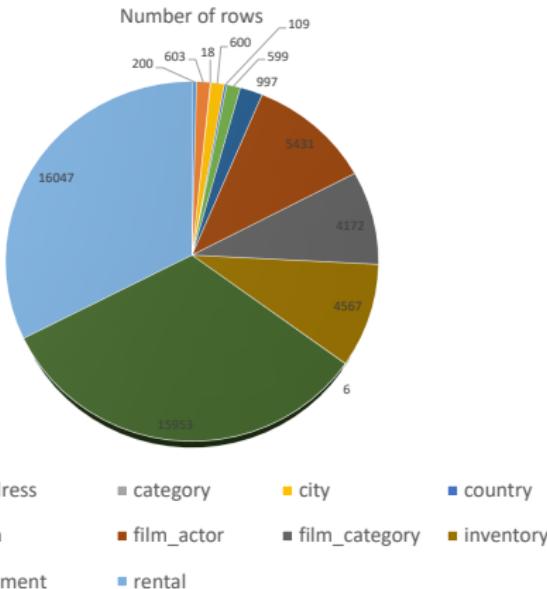
## Lesson 1:

Always check the query plan after an index creation to see whether the index is used or not

## Index Usage

- We need to inspect the query plan (if possible)
- There are two major types of query plan representations:
  - Graphical
  - Text
- Each database system has its own proprietary way how to display a query plan
- The query plan should contain the newly created index name

# Table Statistics of Sakila DB



## Lesson 2:

Do not bother with indexing of small tables

## SARGable Predicate

- ```
SELECT * FROM rental
WHERE YEAR(rental_date) = 2005
    and MONTH(rendal_date) = 10
```
- ```
CREATE INDEX ix_uq ON
rental(rental_date, inventory_id, customer_id)
```

## SARGable Predicate

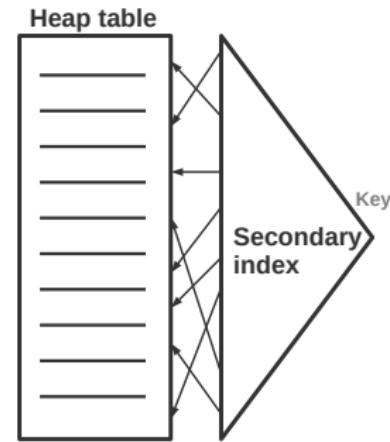
- ```
SELECT * FROM rental
WHERE YEAR(rental_date) = 2005
    and MONTH(rendal_date) = 10
```
- ```
CREATE INDEX ix_uq ON
rental(rental_date, inventory_id, customer_id)
```
- The table is scanned despite the index
- Database systems requires no attribute manipulation in the predicate in order to use an index
- We need to rewrite the query

## Lesson 3:

Do not manipulate the attributes in predicates if possible

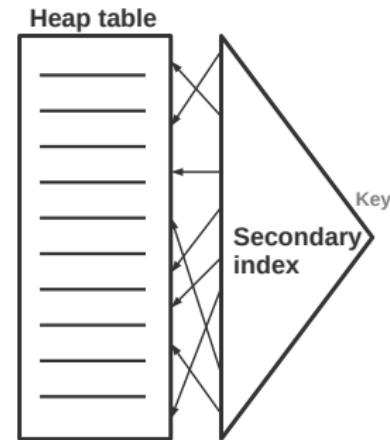
# Asterix

- `SELECT * FROM rental  
WHERE YEAR(rental_date) = 2005  
and MONTH(rendal_date) = 10`
- We may observe that the query plan consists from two data accesses:
  - Secondary index seek
  - Heap table access
- What if we are interested just in few attributes
- `SELECT rental_date, customer_id  
FROM ...`



# Asterix

- ```
SELECT * FROM rental
WHERE YEAR(rental_date) = 2005
    and MONTH(rendal_date) = 10
```
- We may observe that the query plan consists from two data accesses:
  - Secondary index seek
  - Heap table access
- What if we are interested just in few attributes
- ```
SELECT rental_date, customer_id
FROM ...
```



## Lesson 4:

Do not use an asterix in the production

## Covering Indexes

- We avoid table access completely in the previous example
- We use only data in `idx_uq`
- We say that such index *cover* the query
- Having an index with more attributes you are more likely to witness the index to be used during a scan (not during a seek)
- ```
SELECT rental_date, inventory_id, customer_id
FROM rental WHERE inventory_id < 100
```

## Covering Indexes

- We avoid table access completely in the previous example
- We use only data in `idx_uq`
- We say that such index *cover* the query
- Having an index with more attributes you are more likely to witness the index to be used during a scan (not during a seek)
- ```
SELECT rental_date, inventory_id, customer_id
FROM rental WHERE inventory_id < 100
```

# Index Scan vs. Index Seek

```
SELECT rental_date, inventory_id, customer_id
FROM rental
WHERE rental_date >= '2005-10-01'
    and rental_date <= '2005-10-31'
```

rental_date	inventory_id	customer_id
2005-01-02	1	1
2005-01-02	1	50
2005-01-02	200	1
2005-10-02	1	1
2005-10-02	200	1
2005-11-05	99	3
2005-11-05	100	2
2005-11-06	1	1

```
SELECT rental_date, inventory_id, customer_id
FROM rental
WHERE inventory_id < 100
```

rental_date	inventory_id	customer_id
2005-01-02	1	1
2005-01-02	1	50
2005-01-02	200	1
2005-10-02	1	1
2005-10-02	200	1
2005-11-05	99	3
2005-11-05	100	2
2005-11-06	1	1

# Index Scan vs. Index Seek

```
SELECT rental_date, inventory_id, customer_id
FROM rental
WHERE rental_date >= '2005-10-01'
    and rental_date <= '2005-10-31'
```

rental_date	inventory_id	customer_id
2005-01-02	1	1
2005-01-02	1	50
2005-01-02	200	1
2005-10-02	1	1
2005-10-02	200	1
2005-11-05	99	3
2005-11-05	100	2
2005-11-06	1	1

```
SELECT rental_date, inventory_id, customer_id
FROM rental
WHERE inventory_id < 100
```

rental_date	inventory_id	customer_id
2005-01-02	1	1
2005-01-02	1	50
2005-01-02	200	1
2005-10-02	1	1
2005-10-02	200	1
2005-11-05	99	3
2005-11-05	100	2
2005-11-06	1	1

## Insert/Update/Delete and Indexes

- Now lets try an `INSERT` into rental table:

```
INSERT INTO rental
VALUES ('2005-05-26 22:50:00', 367, 130,
        '2005-05-26 22:04:31', 1,
        '2006-02-15 21:30:50')
```

## Insert/Update/Delete and Indexes

- Now lets try an `INSERT` into rental table:

```
INSERT INTO rental
VALUES ('2005-05-26 22:50:00', 367, 130,
        '2005-05-26 22:04:31', 1,
        '2006-02-15 21:30:50')
```

- The database ensures referential integrity during the write operations
- Due to this we should never miss:
  - Indexes on primary keys
  - Indexes on foreign keys

## Lesson 5:

Indexes on PK and FK are important to INSERT, UPDATE and DELETE operations (therefore, not only to SELECT)

# Intelligent Physical Design

- Sometimes it may take a time to come up with intelligent physical design that:
  - supports critical queries
  - do not significantly slow down the insert/update/deletes

# Falkirk Wheel



# Falkirk Wheel



- It lift 550 tons with 22.5 kW

## References

- Course home pages <http://dbedu.cs.vsb.cz>